


ROP

Return Oriented Programming



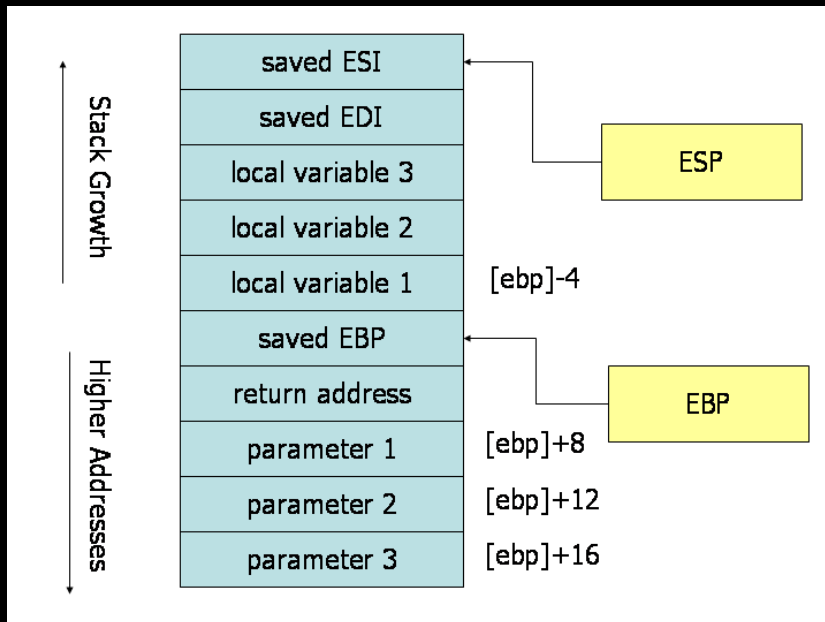
whoami

- briansp8210
- 交通大學資工系大二
- 喜歡學習 pwn 題的相關技巧 XD



Review

- When a program use some dangerous input method like `gets()` or `scanf("%s", buf)` ..., buffer overflow may happen !
- By x86 calling convention, **return address** is stored at **EBP+4**, padding proper length then we can control the program !




[reference](#)

After buffer overflow

Function Epilogue :

```
mov esp, ebp  
pop ebp  
ret
```

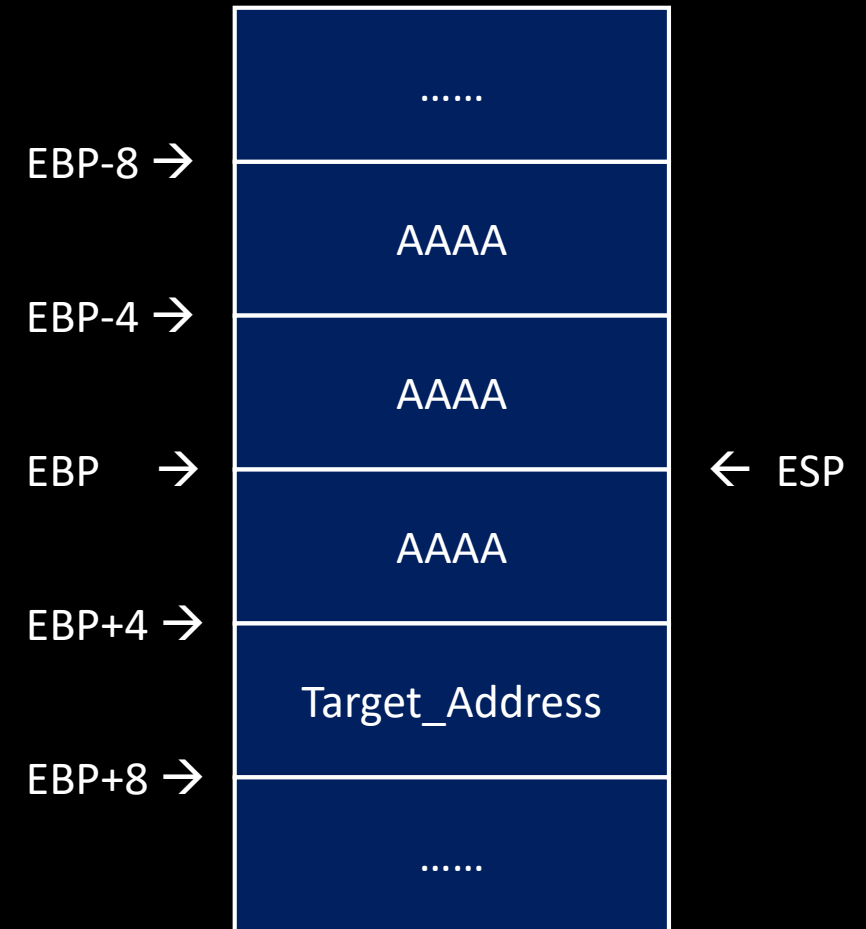


這份投影片裡用黃字標示的指令代表剛跑完這條指令，準備去跑下一條了。

After buffer overflow

Function Epilogue :

```
mov esp, ebp  
pop ebp  
ret
```

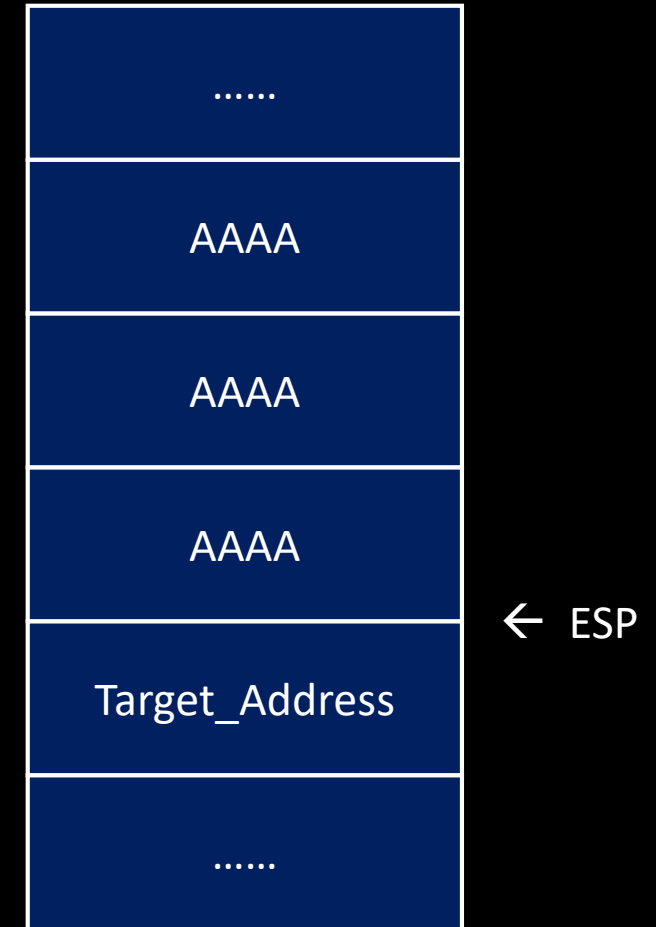


After buffer overflow

Function Epilogue :

```
mov  esp, ebp  
pop  ebp  
ret
```

0x41414141 ← EBP



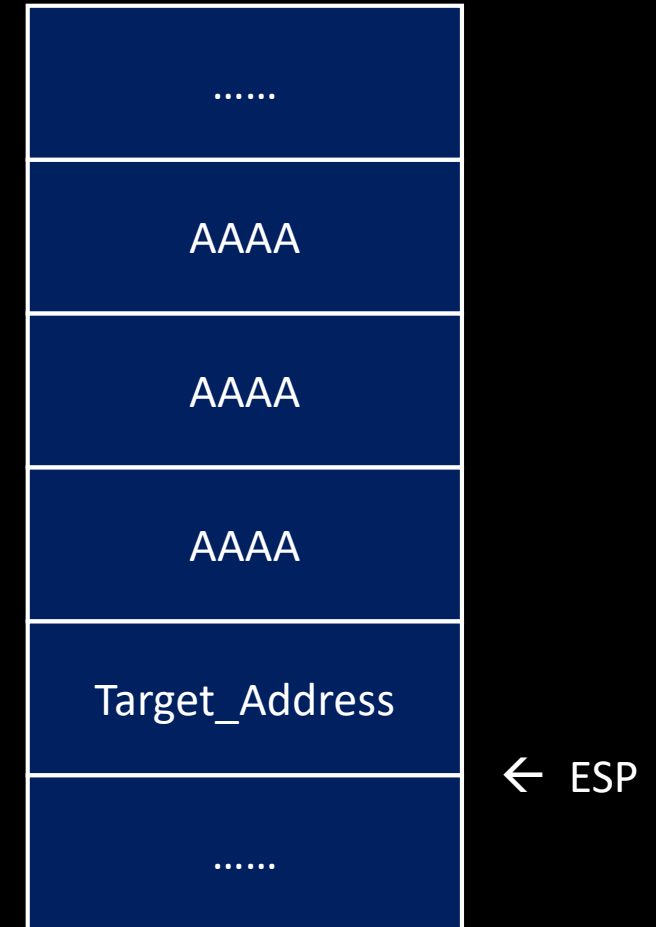
After buffer overflow

Function Epilogue :

```
mov esp, ebp  
pop ebp  
ret
```

0x41414141 ← EBP

Target_Address ← EIP



What is the next step ?

- ret2text
- Return to shellcode
- Return to libc
- ROP

What is the next step ?

- ret2text
- Return to shellcode
- ret2libc
- ROP

ret2text

- When PIE (Position Independent Executable) is enable, program will be loaded to random address, we can't predict where to return.

```
080485fd <secure>:
80485fd: 55          push    ebp
80485fe: 89 e5       mov     ebp,esp
8048600: 83 ec 28    sub     esp,0x28
8048603: c7 04 24 00 00 00 00 mov     DWORD PTR [esp],0x0
804860a: e8 61 fe ff ff call    8048470 <time@plt>
804860f: 89 04 24    mov     DWORD PTR [esp],eax
8048612: e8 99 fe ff ff call    80484b0 <srand@plt>
8048617: e8 c4 fe ff ff call    80484e0 <rand@plt>
804861c: 89 45 f4    mov     DWORD PTR [ebp-0xc],eax
```

```
>>> print ELF('ret2text').checksec()
RELRO: Partial RELRO
Stack: No canary found
NX: NX enabled
PIE: No PIE
```

```
000007eb <secure>:
7eb: 55          push    ebp
7ec: 89 e5       mov     ebp,esp
7ee: 53          push    ebx
7ef: 83 ec 24    sub     esp,0x24
7f2: e8 c9 fe ff ff call    6c0 <__x86.get_pc_thunk.bx>
7f7: 81 c3 09 18 00 00 add     ebx,0x1809
7fd: c7 04 24 00 00 00 00 mov     DWORD PTR [esp],0x0
804: e8 d7 fd ff ff call    5e0 <time@plt>
809: 89 04 24    mov     DWORD PTR [esp],eax
80c: e8 1f fe ff ff call    630 <srand@plt>
811: e8 4a fe ff ff call    660 <rand@plt>
816: 89 45 f4    mov     DWORD PTR [ebp-0xc],eax
```

```
>>> print ELF('ret2text_with_PIE').checksec()
RELRO: Partial RELRO
Stack: No canary found
NX: NX enabled
PIE: PIE enabled
```

- If not, we can use desirable program code to do exploit !

ret2text

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

void secure(void)
{
    int secretcode, input;
    srand(time(NULL));

    secretcode = rand();
    scanf("%d", &input);
    if(input == secretcode)
        system("/bin/sh");
}

int main(void)
{
    setvbuf(stdout, 0LL, 2, 0LL);
    setvbuf(stdin, 0LL, 1, 0LL);

    char buf[100];

    printf("There is something amazing here, do you know anything?\n");
    gets(buf);
    printf("Maybe I will tell you next time !");

    return 0;
}
```

What is the next step ?

- ret2text
- Return to shellcode
- ret2libc
- ROP

Return to shellcode

- When memory address isn't marked as non-executable, we can send shellcode to it and return to its address to run our shellcode.

We can both
write and
execute here !

| | | | | | |
|-------------------|------|----------|-------|---------|---|
| 08048000-08049000 | r-xp | 00000000 | 08:01 | 928493 | /root/CSC/ROP/return_to_shellcode/return_to_shellcode |
| 08049000-0804a000 | r-xp | 00000000 | 08:01 | 928493 | /root/CSC/ROP/return_to_shellcode/return_to_shellcode |
| 0804a000-0804b000 | rwpx | 00001000 | 08:01 | 928493 | /root/CSC/ROP/return_to_shellcode/return_to_shellcode |
| 09df9000-09e1a000 | rwpx | 00000000 | 00:00 | 0 | [heap] |
| f7599000-f759a000 | rwpx | 00000000 | 00:00 | 0 | |
| f759a000-f7746000 | r-xp | 00000000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7746000-f7747000 | ---p | 001ac000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7747000-f7749000 | r-xp | 001ac000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7749000-f774a000 | rwpx | 001ae000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f774a000-f774e000 | rwpx | 00000000 | 00:00 | 0 | |
| f776d000-f776f000 | r--p | 00000000 | 00:00 | 0 | [vvar] |
| f776f000-f7770000 | r-xp | 00000000 | 00:00 | 0 | [vdso] |
| f7770000-f7792000 | r-xp | 00000000 | 08:01 | 3014701 | /lib32/ld-2.23.so |
| f7792000-f7793000 | rwpx | 00000000 | 00:00 | 0 | |
| f7793000-f7794000 | r-xp | 00022000 | 08:01 | 3014701 | /lib32/ld-2.23.so |
| f7794000-f7795000 | rwpx | 00023000 | 08:01 | 3014701 | /lib32/ld-2.23.so |
| ffd96000-ffdb7000 | rwpx | 00000000 | 00:00 | 0 | [stack] |

```
>>> print ELF('return_to_shellcode').checksec()
RELRO: Partial RELRO
Stack: No canary found
NX: NX disabled
PIE: No PIE
```

Return to shellcode

```
#include <stdio.h>
#include <string.h>

char buf2[100];

int main(void)
{
    setvbuf(stdout, 0LL, 2, 0LL);
    setvbuf(stdin, 0LL, 1, 0LL);

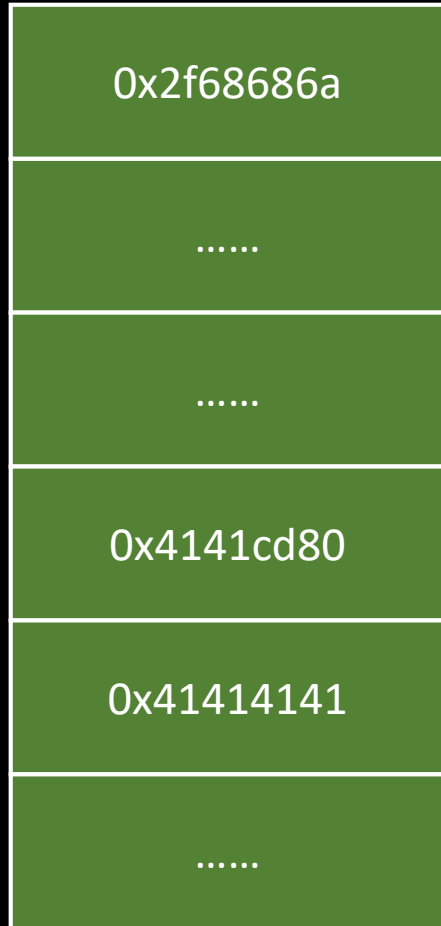
    char buf[100];

    printf("No system for you this time !!!\n");
    gets(buf);
    strncpy(buf2, buf, 100);
    printf("bye bye ~");

    return 0;
}
```


Return to shellcode

buf2 →



data buffer

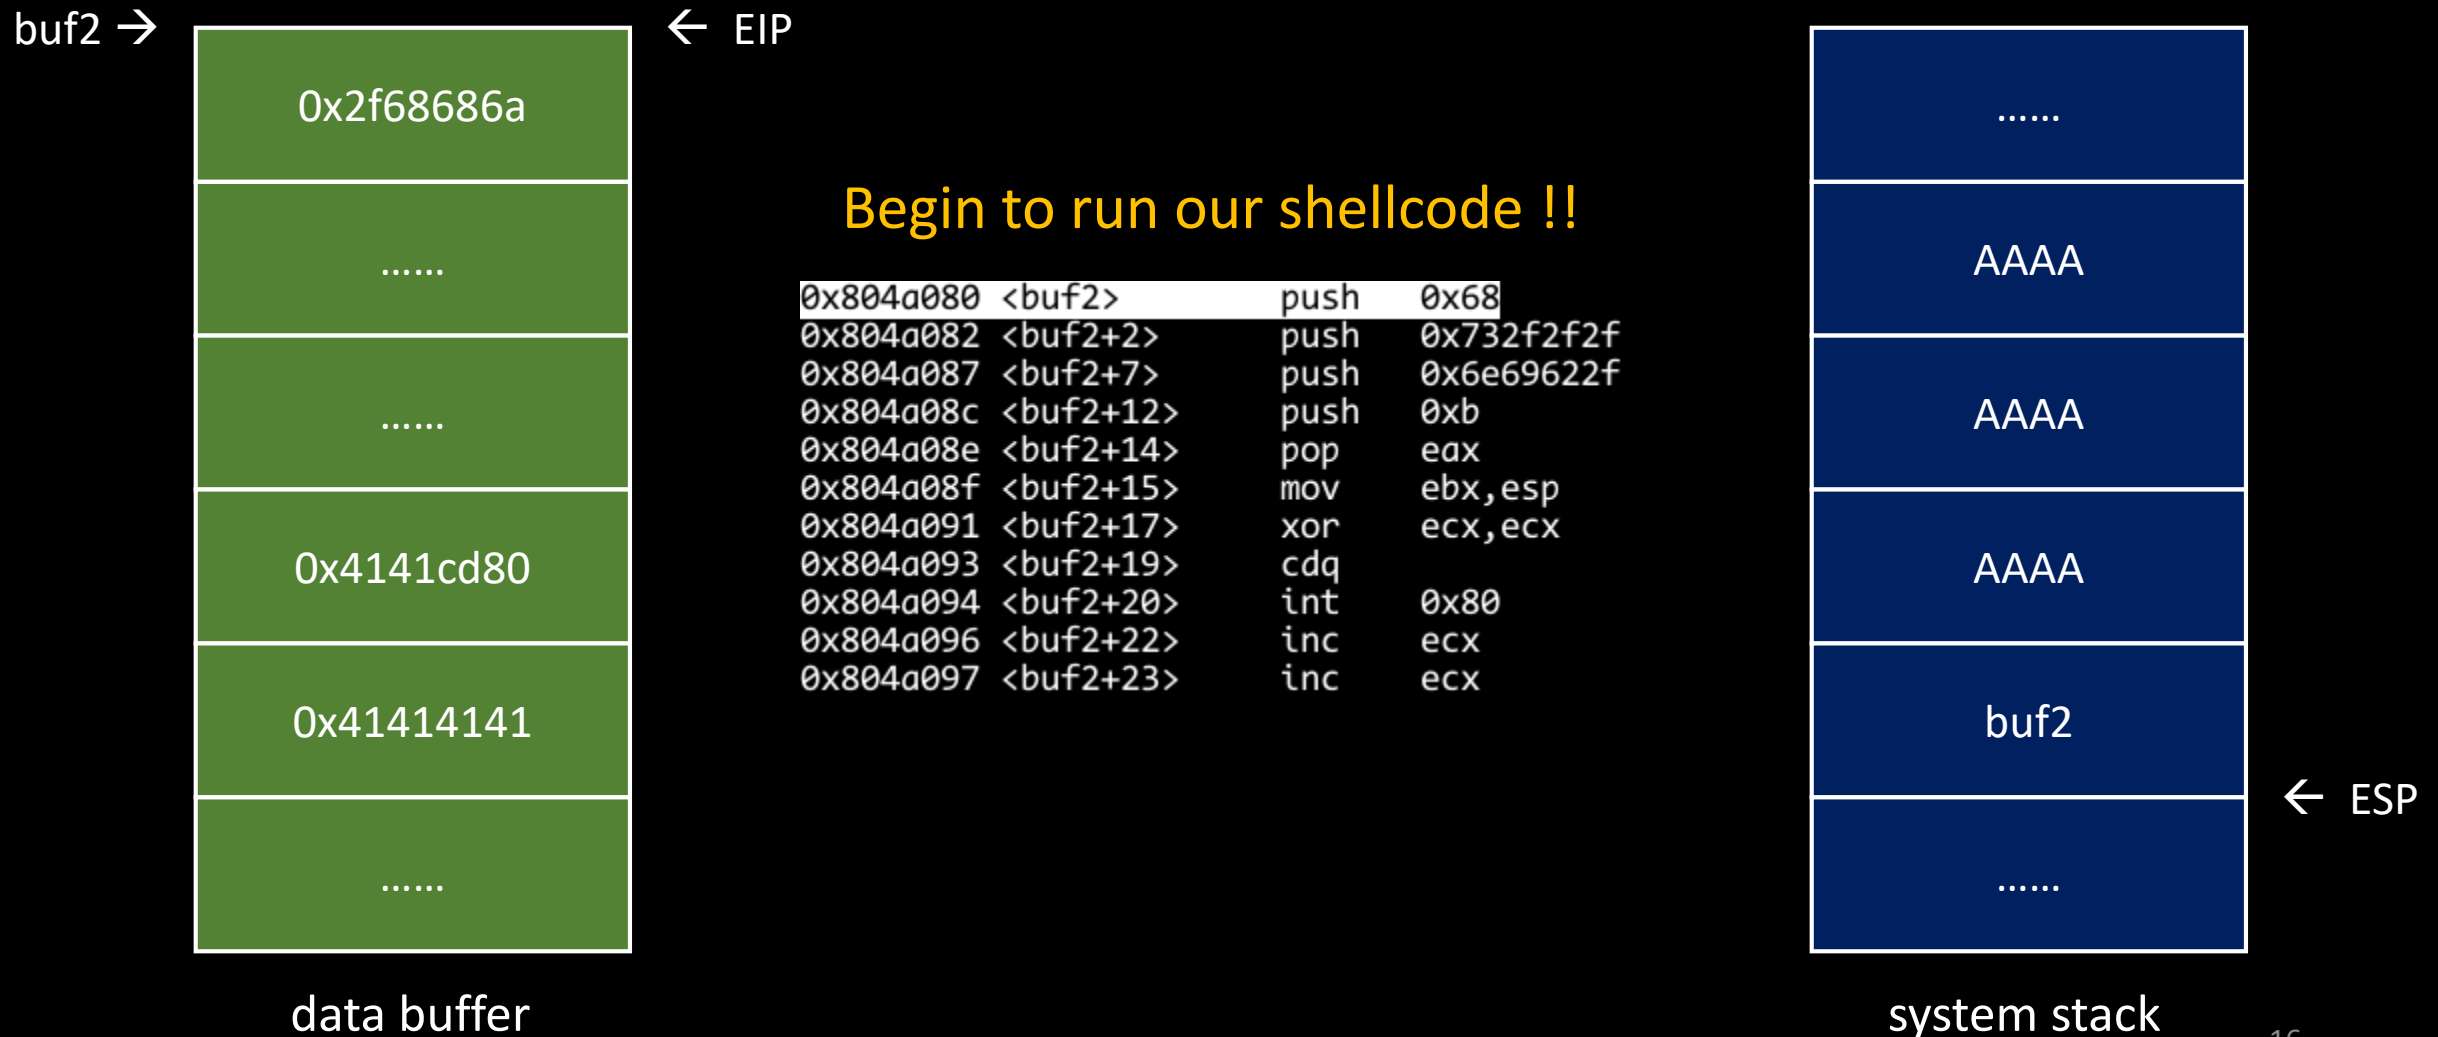
```
0x80485a8 <main+123>    mov     DWORD PTR [esp],0x804a080
0x80485af <main+130>    call    0x8048420 <strncpy@plt>
0x80485b4 <main+135>    mov     DWORD PTR [esp],0x8048680
0x80485bb <main+142>    call    0x80483c0 <printf@plt>
0x80485c0 <main+147>    mov     eax,0x0
0x80485c5 <main+152>    leave
0x80485c6 <main+153>    ret
```



← ESP

system stack

Return to shellcode



What is the next step ?

- ret2text
- Return to shellcode
- ret2libc
- ROP

Call a function

- Observe normal function call
`strcpy (Destination, Source) ;`

```
.....  
mov     DWORD PTR [esp+0x4],0x804a020  
mov     DWORD PTR [esp],0x804a035  
call    80482f0 <strcpy@plt>  
.....
```

Source address: 0x804a020

Destination address: 0x804a035




Call a function

- Observe normal function call
`strcpy(Destination, Source) ;`

```
.....  
mov     DWORD PTR [esp+0x4],0x804a020  
mov     DWORD PTR [esp],0x804a035  
call    80482f0 <strcpy@plt>  
.....
```

Source address: 0x804a020

Destination address: 0x804a035



Call a function

- Observe normal function call
`strcpy(Destination, Source) ;`

```
.....  
mov     DWORD PTR [esp+0x4],0x804a020  
mov     DWORD PTR [esp],0x804a035  
call    80482f0 <strcpy@plt>  
.....
```

Source address: 0x804a020

Destination address: 0x804a035



Call a function

- Observe normal function call
`strcpy(Destination, Source) ;`

```
.....  
mov     DWORD PTR [esp+0x4],0x804a020  
mov     DWORD PTR [esp],0x804a035  
call    80482f0 <strcpy@plt>  
.....
```

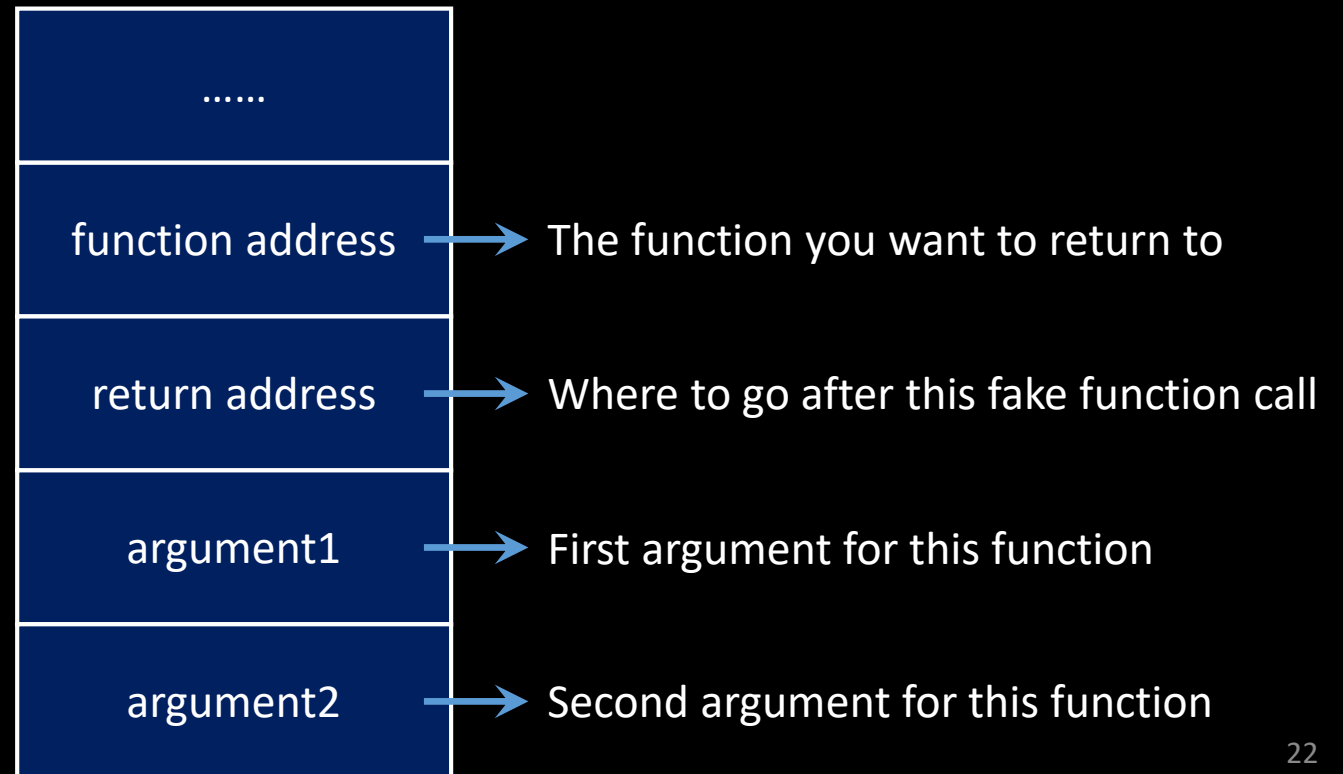
Source address: 0x804a020

Destination address: 0x804a035




ret2libc

- When **calling** a function, there will be some needed information on stack.
- If we know a function's address, we can fake a function call by **return** to the function with proper value on stack.



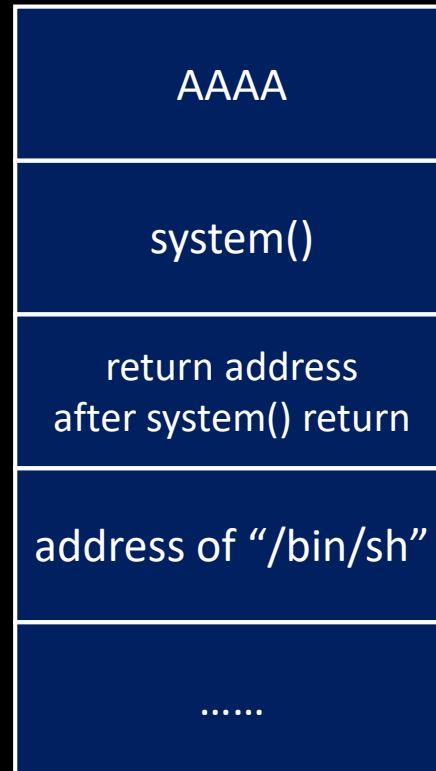
fake a function call

- Put needed value on stack when buffer overflow .
- After returning to the function, it can fetch right data from right position, just like a normal function call !



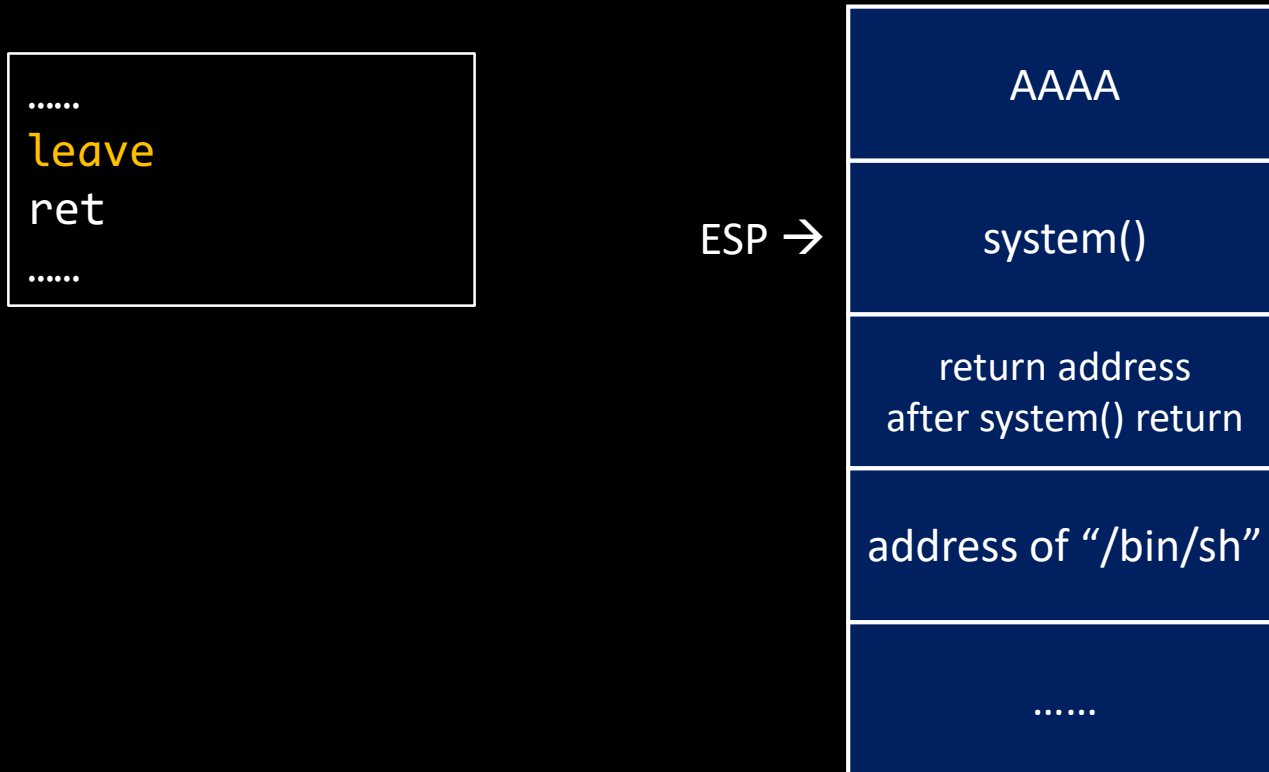
fake a function call

- Put needed value on stack when buffer overflow .
- After returning to the function, it can fetch right data from right position, just like a normal function call !



fake a function call

- Put needed value on stack when buffer overflow .
- After returning to the function, it can fetch right data from right position, just like a normal function call !

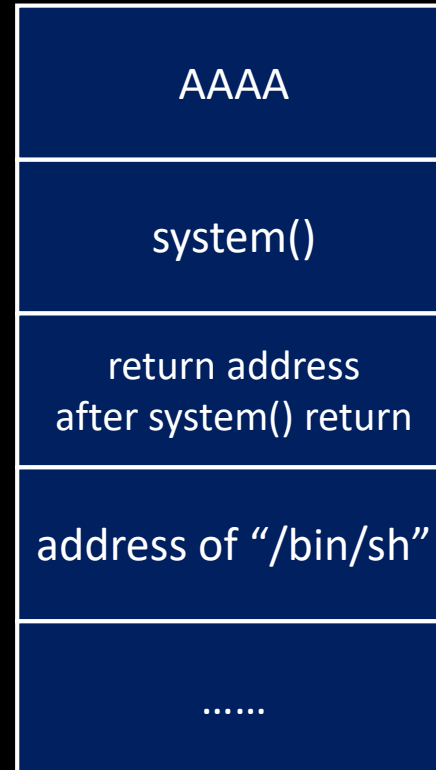


fake a function call

- Put needed value on stack when buffer overflow .
- After returning to the function, it can fetch right data from right position, just like a normal function call !

```
.....  
leave  
ret  
.....
```

ESP →



ret2libc

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

char *shell = "/bin/sh";
char buf2[100];

void secure(void)
{
    int secretcode, input;
    srand(time(NULL));

    secretcode = rand();
    scanf("%d", &input);
    if(input == secretcode)
        system("shell!?");
}

int main(void)
{
    setvbuf(stdout, 0LL, 2, 0LL);
    setvbuf(stdin, 0LL, 1, 0LL);

    char buf1[100];

    printf("RET2LIBC >_<\n");
    gets(buf1);

    return 0;
}
```

ret2libc

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

char buf2[100];

void secure(void)
{
    int secretcode, input;
    srand(time(NULL));

    secretcode = rand();
    scanf("%d", &input);
    if(input == secretcode)
        system("no_shell_QQ");
}

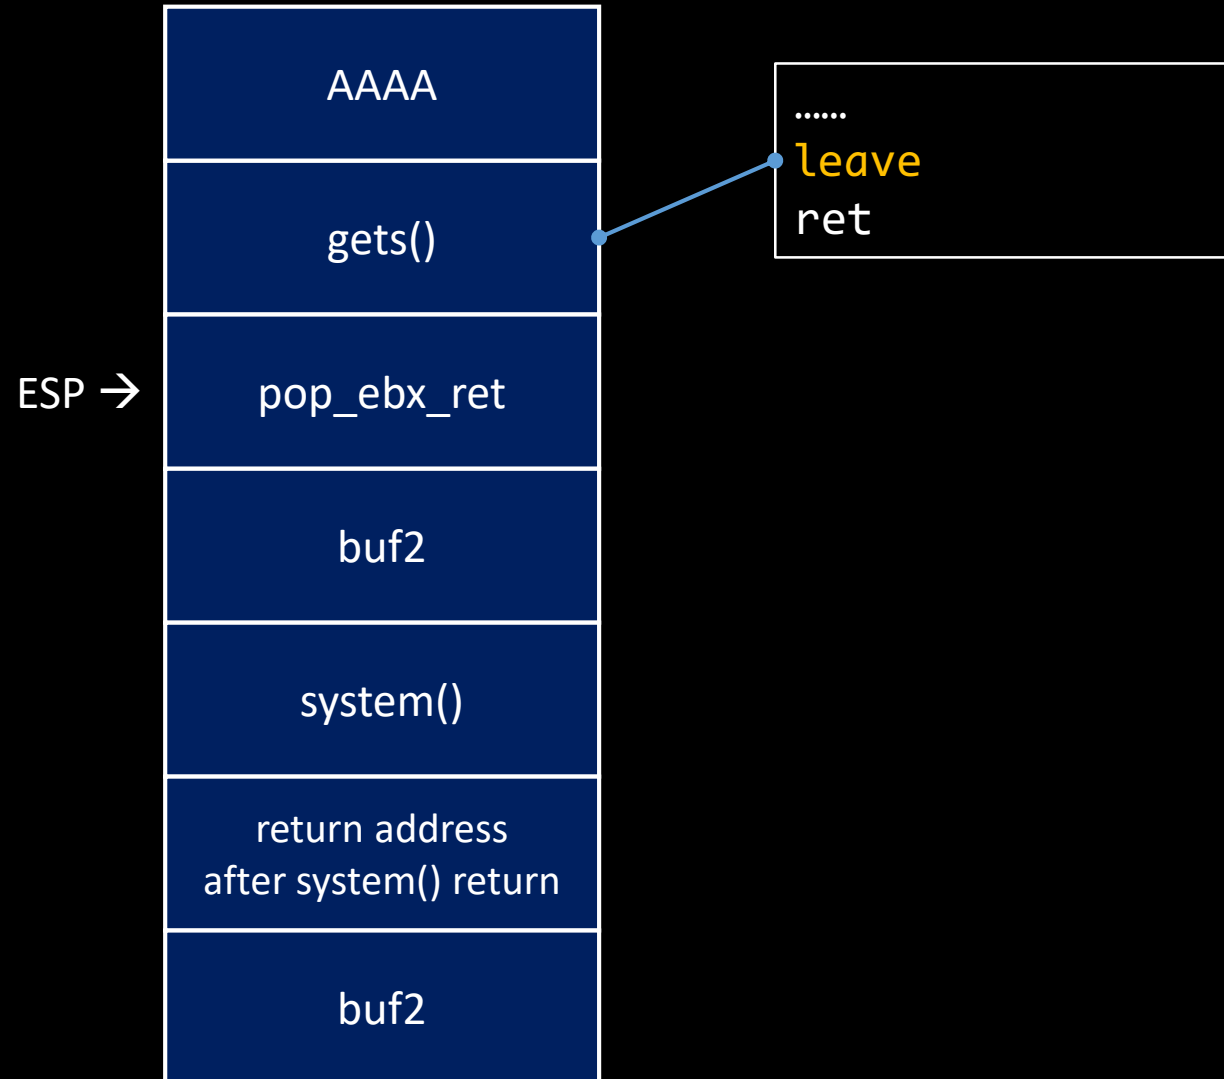
int main(void)
{
    setvbuf(stdout, 0LL, 2, 0LL);
    setvbuf(stdin, 0LL, 1, 0LL);

    char buf1[100];

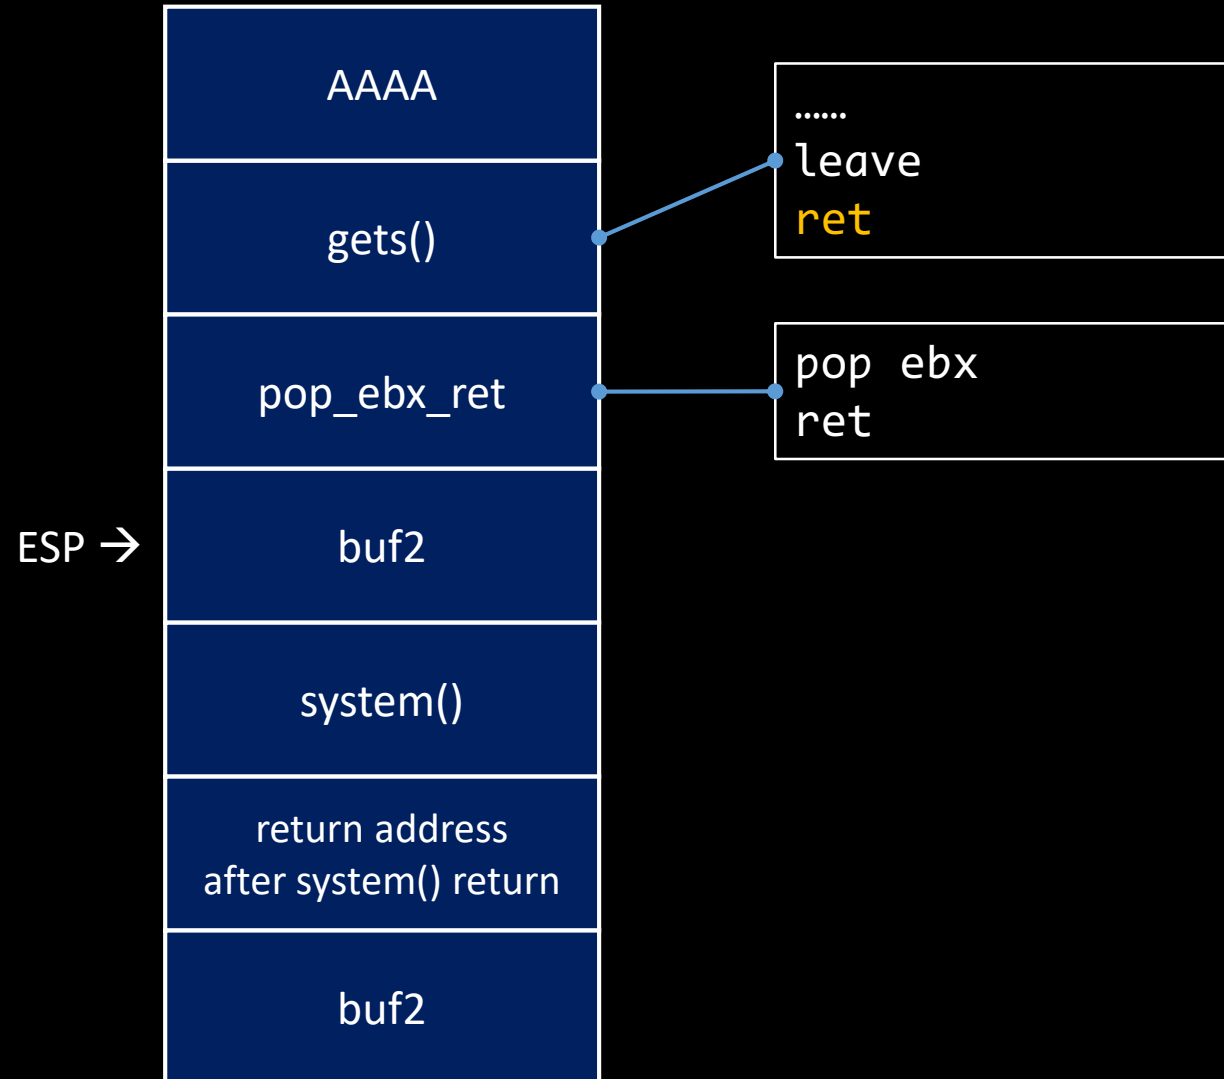
    printf("Something surprise here, but I don't think it will work.\n");
    printf("What do you think ?");
    gets(buf1);

    return 0;
}
```

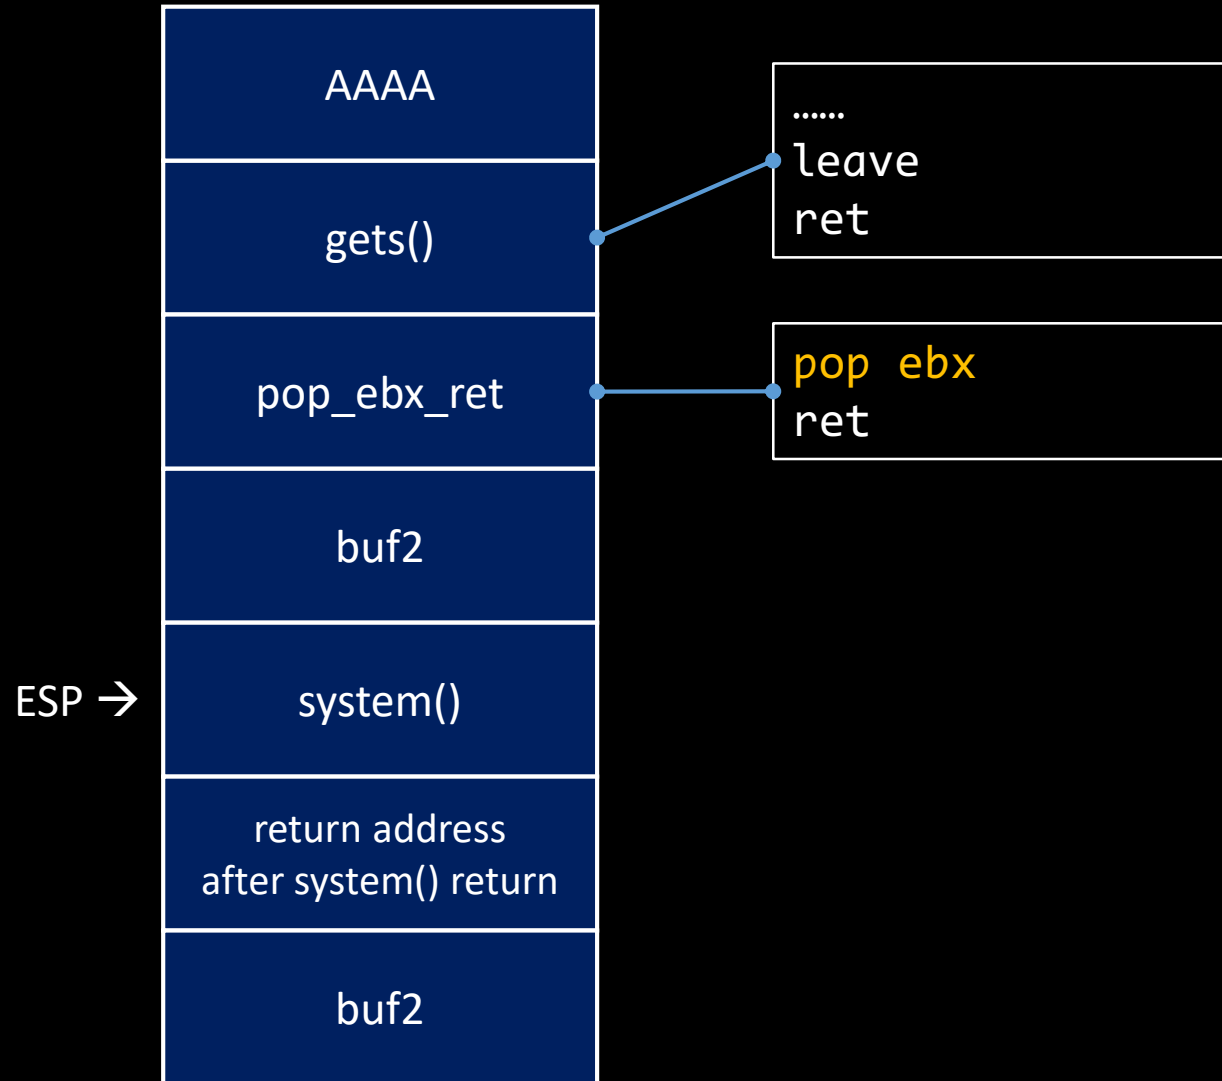
ret2libc



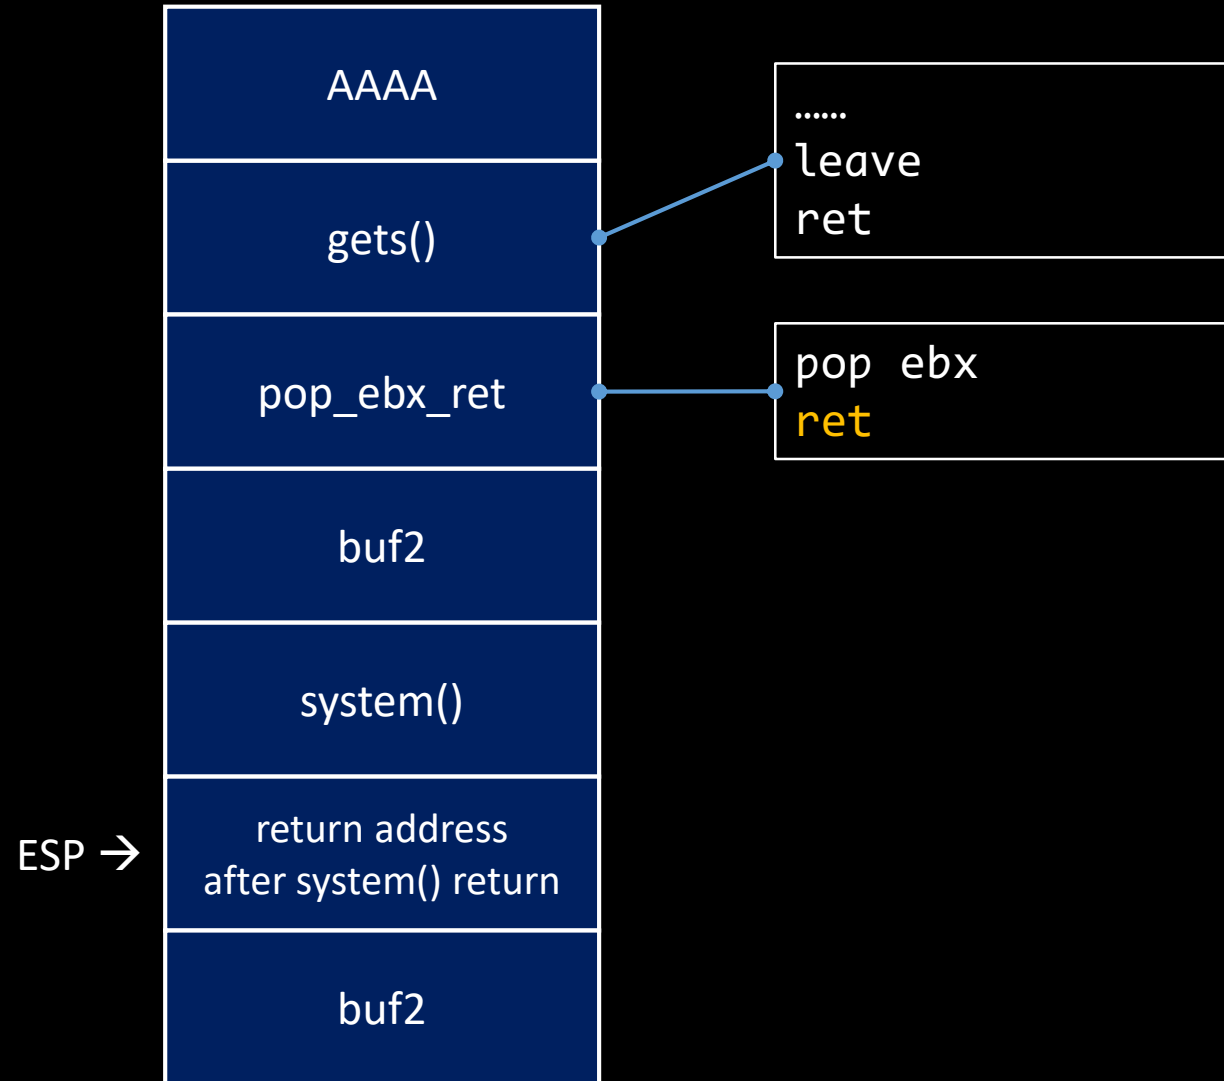
ret2libc



ret2libc



ret2libc



ASLR and libc address

- With ASLR, the base address of libc is randomized each time we execute a program .
- We can't predict the function's address we want to use and return to it, except those used by this program .

| | | | | | |
|-------------------|------|----------|-------|---------|------------------------|
| 08048000-08049000 | r-xp | 00000000 | 08:01 | 928487 | /root/CSC/ROP/ret2text |
| 08049000-0804a000 | r--p | 00000000 | 08:01 | 928487 | /root/CSC/ROP/ret2text |
| 0804a000-0804b000 | rw-p | 00001000 | 08:01 | 928487 | /root/CSC/ROP/ret2text |
| 08dc2000-08de3000 | rw-p | 00000000 | 00:00 | 0 | [heap] |
| f75a6000-f75a7000 | rw-p | 00000000 | 00:00 | 0 | |
| f75a7000-f7753000 | r-xp | 00000000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7753000-f7754000 | ---p | 001ac000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7754000-f7756000 | r--p | 001ac000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7756000-f7757000 | rw-p | 001ae000 | 08:01 | 3014705 | /lib32/libc-2.23.so |
| f7757000-f775b000 | rw-p | 00000000 | 00:00 | 0 | |
| f777a000-f777c000 | r--p | 00000000 | 00:00 | 0 | [vvar] |
| f777c000-f777d000 | r-xp | 00000000 | 00:00 | 0 | [vdso] |
| f777d000-f779f000 | r-xp | 00000000 | 08:01 | 3014701 | /lib32/ld-2.23.so |
| f779f000-f77a0000 | rw-p | 00000000 | 00:00 | 0 | |
| f77a0000-f77a1000 | r--p | 00022000 | 08:01 | 3014701 | /lib32/ld-2.23.so |
| f77a1000-f77a2000 | rw-p | 00023000 | 08:01 | 3014701 | /lib32/ld-2.23.so |
| ffa90000-ffab1000 | rw-p | 00000000 | 00:00 | 0 | [stack] |

How to find functions' offset

- The offset of a function in a libc will be fixed .
- Function's address will be **libc base + function's offset** .
- readelf or pwntools can help you find the offset !

| | | | |
|---------------|----------|----------------|--------------------------------|
| 192: 001288b0 | 91 FUNC | GLOBAL DEFAULT | 12 getnetbyname_r@GLIBC_2.0 |
| 193: 0007a400 | 82 IFUNC | GLOBAL DEFAULT | 12 strcmp@@GLIBC_2.0 |
| 194: 000fd0e0 | 607 FUNC | GLOBAL DEFAULT | 12 getnetbyname_r@@GLIBC_2.1.2 |
| 195: 000fd860 | 211 FUNC | GLOBAL DEFAULT | 12 getprotoent_r@@GLIBC_2.1.2 |
| 196: 001175a0 | 17 FUNC | GLOBAL DEFAULT | 12 svcfd_create@@GLIBC_2.0 |
| 197: 000e3d20 | 58 FUNC | WEAK DEFAULT | 12 ftruncate@@GLIBC_2.0 |
| 198: 00128970 | 53 FUNC | GLOBAL DEFAULT | 12 getprotoent_r@GLIBC_2.0 |
| 199: 00081150 | 44 FUNC | GLOBAL DEFAULT | 12 __strncpy_gg@@GLIBC_2.1.1 |
| 200: 0010f4c0 | 138 FUNC | GLOBAL DEFAULT | 12 xdr_unixcred@@GLIBC_2.1 |
| 201: 00029730 | 72 FUNC | WEAK DEFAULT | 12 dcngettext@@GLIBC_2.2 |
| 202: 0010cd50 | 125 FUNC | GLOBAL DEFAULT | 12 xdr_rmtcallres@@GLIBC_2.0 |
| 203: 00064da0 | 421 FUNC | GLOBAL DEFAULT | 12 _IO_puts@@GLIBC_2.0 |
| 204: 001075d0 | 242 FUNC | GLOBAL DEFAULT | 12 inet_nsap_addr@@GLIBC_2.0 |
| 205: 00106de0 | 286 FUNC | WEAK DEFAULT | 12 inet_aton@@GLIBC_2.0 |
| 206: 000e4890 | 217 FUNC | GLOBAL DEFAULT | 12 ttyslot@@GLIBC_2.0 |
| 207: 001aa8dc | 4 OBJECT | GLOBAL DEFAULT | 32 __rcmd_errstr@@GLIBC_2.0 |
| 208: 000d7dd0 | 90 FUNC | GLOBAL DEFAULT | 12 wordfree@@GLIBC_2.1 |

`readelf -s file_name`

```
>>> from pwn import *
>>> libc = ELF('libc.so.6')
[*] '/root/CSC/ROP/libc.so.6'
Arch:      i386-32-little
RELRO:     Partial RELRO
Stack:     Canary found
NX:        NX enabled
PIE:       PIE enabled
>>> print hex(libc.symbols['strcmp'])
0x7a400
>>>
```

How to find libc's base address

- For those dynamically linked program, when a function is called once, its address in libc will be written to its GOT entry .

file *file_name*


ret2text: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), dynamically linked, interpreter /lib/ld-linux.so.2, for GNU/Linux 2.6.24, BuildID[!]

objdump -d *file_name*

```
080482f0 <strcpy@plt>:
80482f0: ff 25 0c a0 04 08    jmp     DWORD PTR ds:0x804a00c
80482f6: 68 00 00 00 00      push    0x0
80482fb: e9 e0 ff ff ff      jmp     80482e0 <_init+0x2c>

08048300 <__gmon_start__@plt>:
8048300: ff 25 10 a0 04 08    jmp     DWORD PTR ds:0x804a010
8048306: 68 08 00 00 00      push    0x8
804830b: e9 d0 ff ff ff      jmp     80482e0 <_init+0x2c>

08048310 <__libc_start_main@plt>:
8048310: ff 25 14 a0 04 08    jmp     DWORD PTR ds:0x804a014
8048316: 68 10 00 00 00      push    0x10
804831b: e9 c0 ff ff ff      jmp     80482e0 <_init+0x2c>
```

 <strcpy@got.plt> , this entry will store :

next instruction in <strcpy@plt>
if strcpy haven't been called yet .

strcpy()'s address in libc !
if strcpy have been called .

How to find libc's base address

- For those dynamically linked program, when a function is called once, its address in libc will be written to its GOT entry .

file *file_name*

ret2text: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), dynamically linked, interpreter /lib/ld-linux.so.2, for GNU/Linux 2.6.24, BuildID[!]

objdump -d *file_name*

```
080482f0 <strcpy@plt>:
80482f0:    ff 25 0c a0 04 08    jmp     DWORD PTR ds:0x804a00c
80482f6:    68 00 00 00 00      push    0x0
80482fb:    e9 e0 ff ff ff      jmp     80482e0 <_init+0x2c>

08048300 <__gmon_start__@plt>:
8048300:    ff 25 10 a0 04 08    jmp     DWORD PTR ds:0x804a010
8048306:    68 08 00 00 00      push    0x8
804830b:    e9 d0 ff ff ff      jmp     80482e0 <_init+0x2c>

08048310 <__libc_start_main@plt>:
8048310:    ff 25 14 a0 04 08    jmp     DWORD PTR ds:0x804a014
8048316:    68 10 00 00 00      push    0x10
804831b:    e9 c0 ff ff ff      jmp     80482e0 <_init+0x2c>
```

<strcpy@got.plt> , this entry will store :

next instruction in <strcpy@plt>
if strcpy haven't been called yet .

strcpy()'s address in libc !
if strcpy have been called .

How to find libc's base address

- We can use `puts` to leak a function's address store in GOT entry .
- Subtract this function's offset in libc from the leaked address, we can get libc's base address !
- Then we can know each function's address :
libc base + function's offsets

ret2libc

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

char buf2[100];

void secure(void)
{
    int secretcode, input;
    srand(time(NULL));

    secretcode = rand();
    scanf("%d", &input);
    if(input == secretcode)
        puts("no_shell_QQ");
}

int main(void)
{
    setvbuf(stdout, 0LL, 2, 0LL);
    setvbuf(stdin, 0LL, 1, 0LL);

    char buf1[100];

    printf("No surprise anymore, system disappeared QQ.\n");
    printf("Can you find it !?");
    gets(buf1);

    return 0;
}
```

What is the next step ?

- ret2text
- Return to shellcode
- ret2libc
- ROP


ROP

- Try to use gadgets, some machine instruction sequences which end with `ret`, to achieve our goal .

| | | | |
|----------|----------|-----|----------|
| 8048729: | 83 c4 1c | add | esp,0x1c |
| 804872c: | 5b | pop | ebx |
| 804872d: | 5e | pop | esi |
| 804872e: | 5f | pop | edi |
| 804872f: | 5d | pop | ebp |
| 8048730: | c3 | ret | |

For example, this is a gadget comes from `<__libc_csu_init>`

ROP



ROP

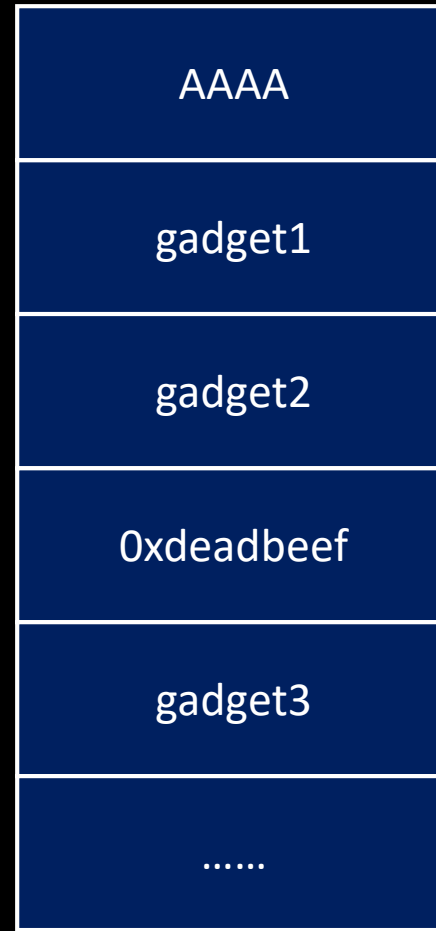
| |
|------------|
| AAAA |
| gadget1 |
| gadget2 |
| 0xdeadbeef |
| gadget3 |
| |



ROP

```
.....  
leave  
ret  
.....
```

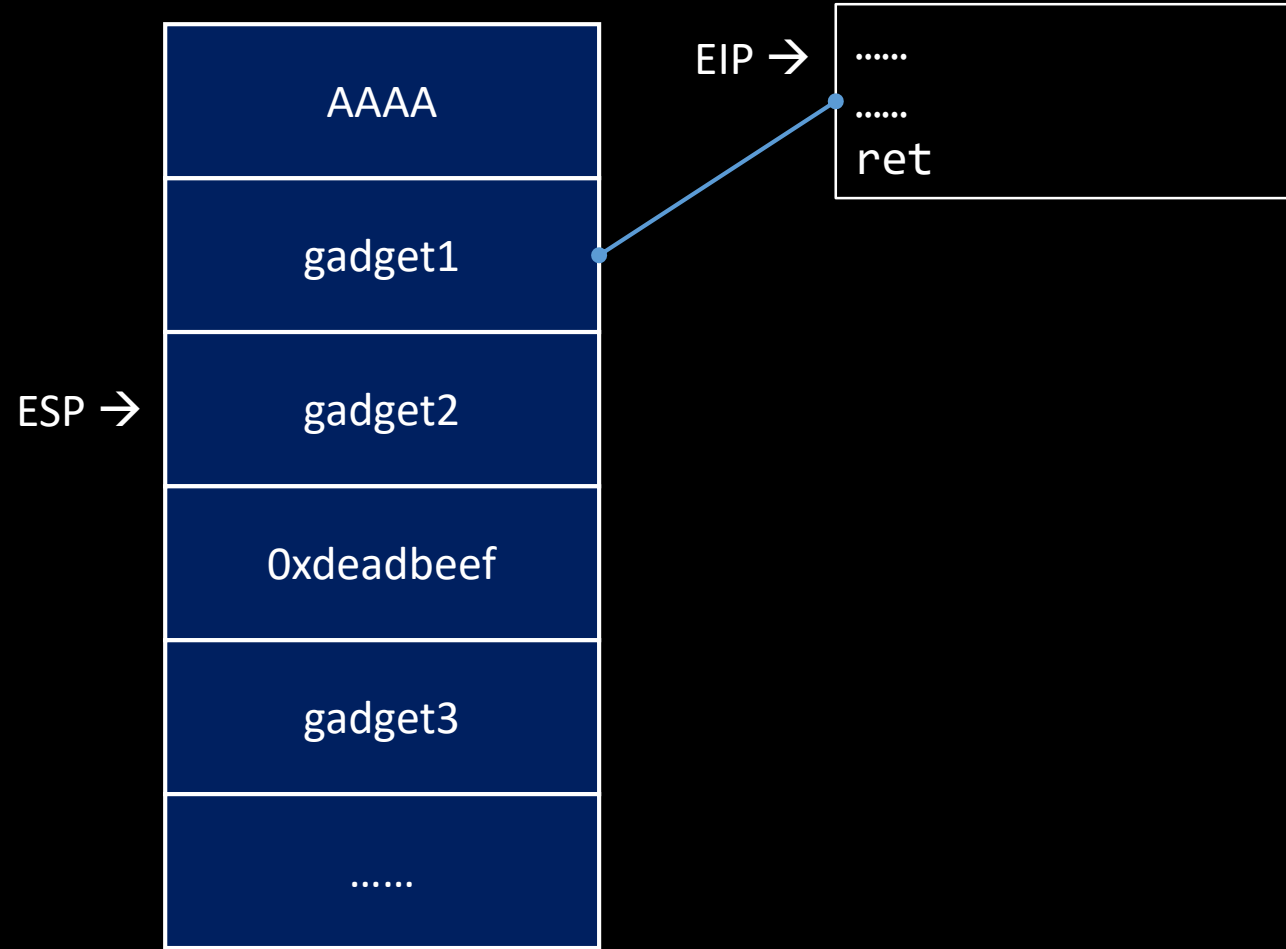
ESP →



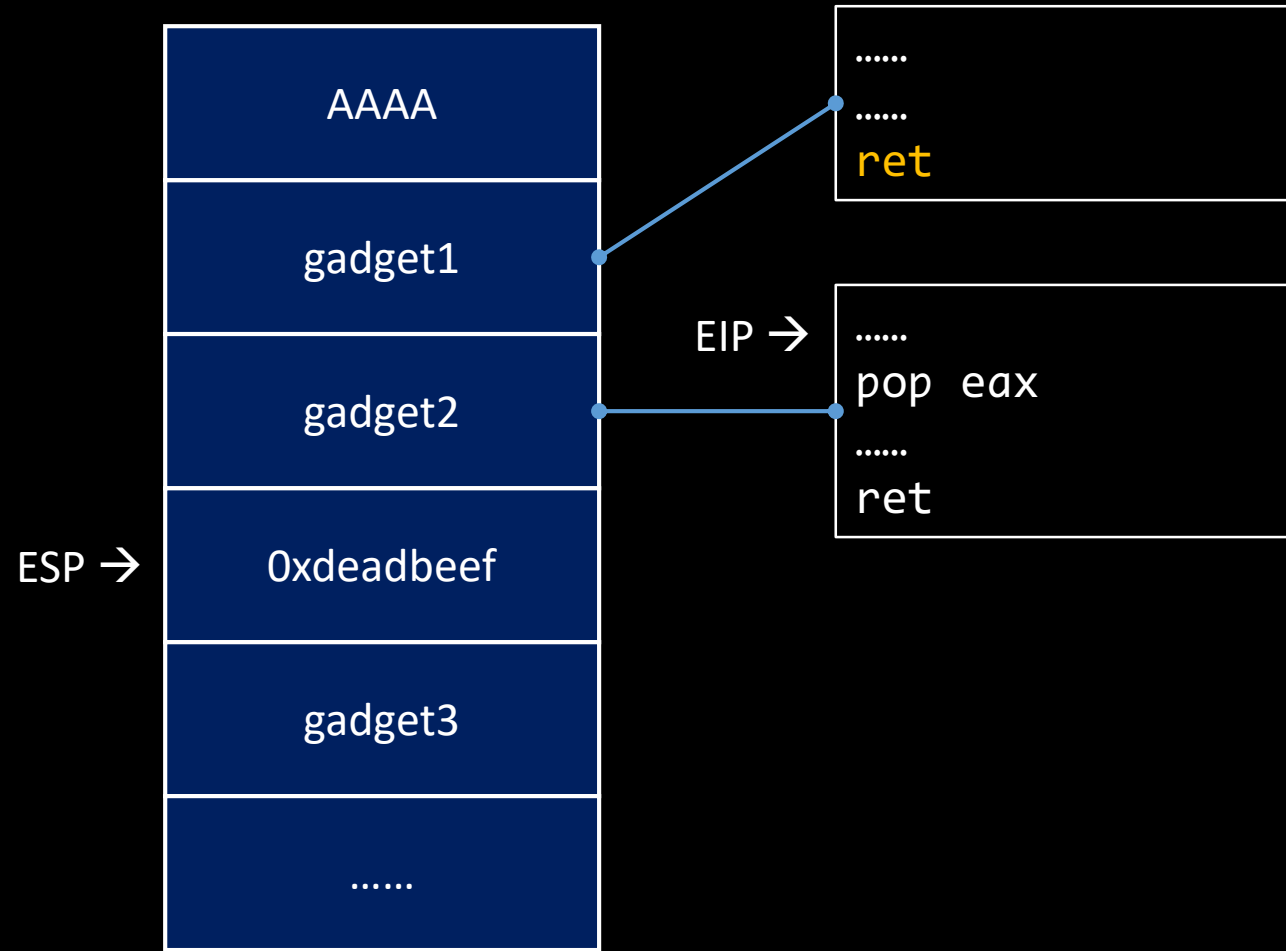
buffer overflow !!

ROP

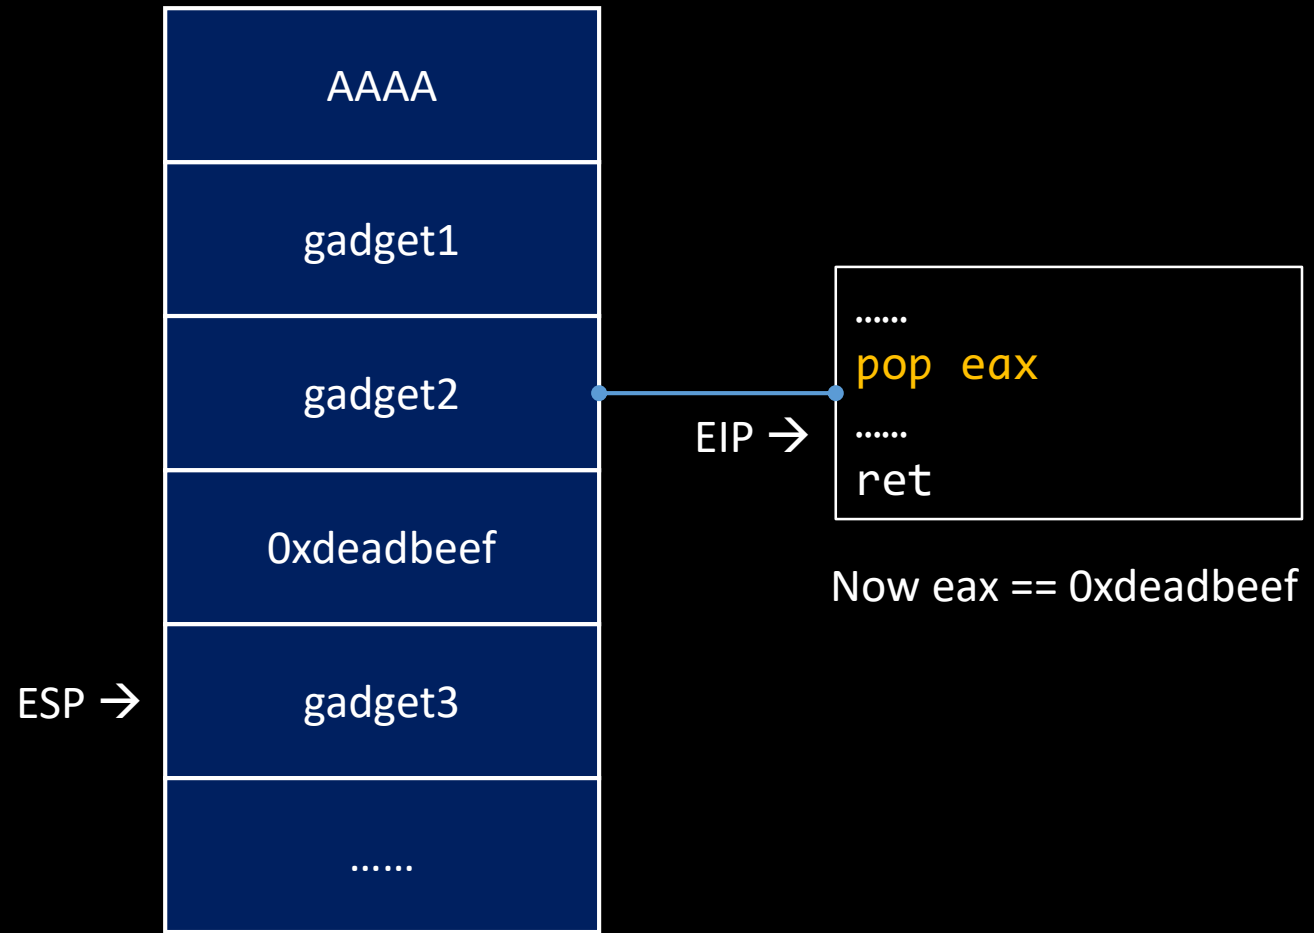
```
.....  
leave  
ret  
.....
```



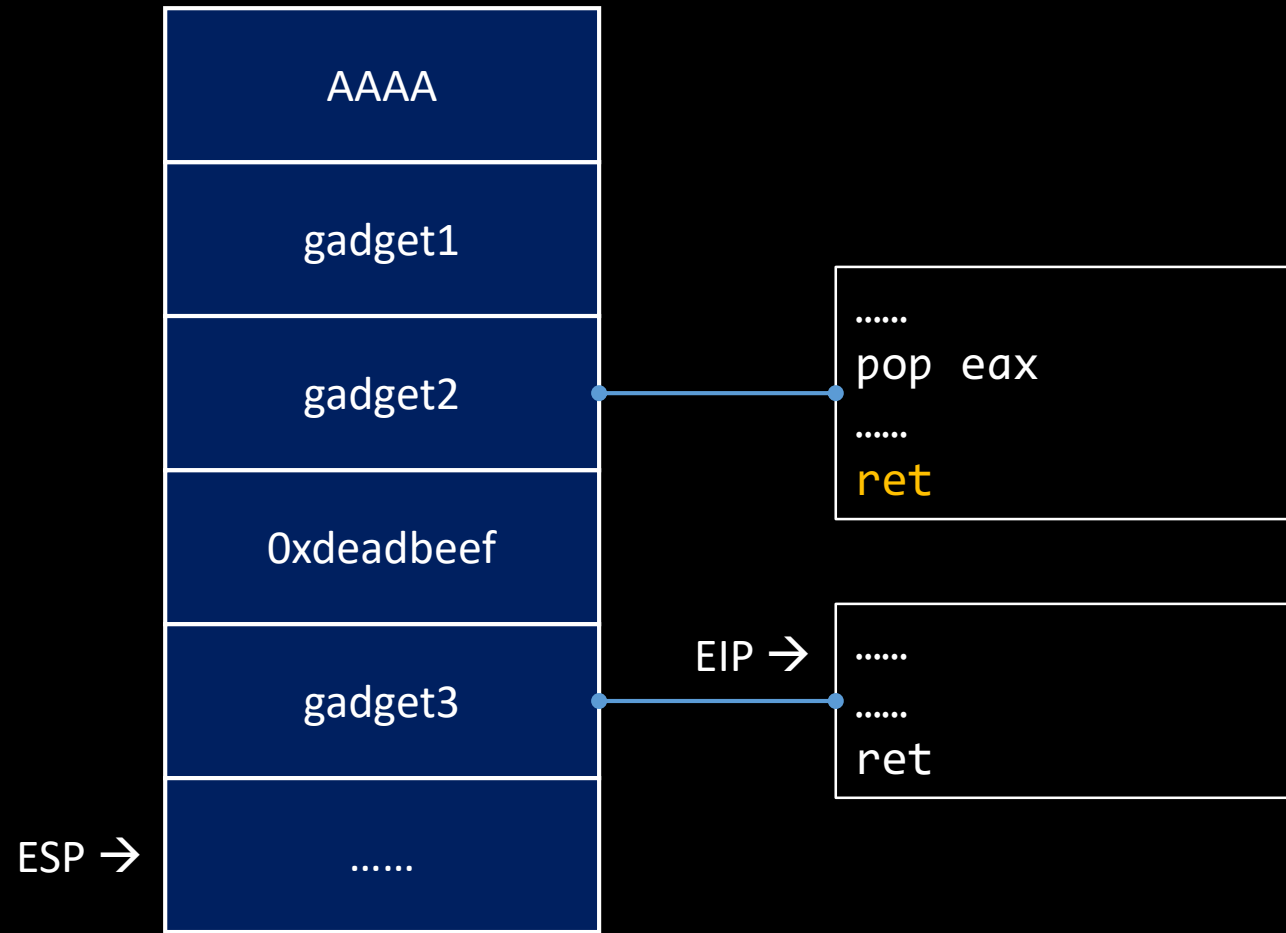
ROP



ROP



ROP



Find available gadgets

- ROPgadget can help you find gadgets .

- Usage:

ROPgadget --binary *file_name*

ROPgadget --binary *file_name* --opcode *opcode*

Gadgets information

```
=====
0x080488a3 : adc al, 0x41 ; ret
0x08048560 : add al, 0x24 ; inc eax ; mov al, byte ptr [0xd0ff0804] ; leave ; ret
0x0804859d : add al, 0x24 ; inc eax ; mov al, byte ptr [0xd2ff0804] ; leave ; ret
0x080485c8 : add al, 8 ; add ecx, ecx ; ret
0x08048564 : add al, 8 ; call eax
0x080485a1 : add al, 8 ; call edx
0x08048548 : add al, 8 ; cmp eax, 6 ; ja 0x8048557 ; ret
0x080486c0 : add byte ptr [eax], al ; add byte ptr [eax], al ; leave ; ret
0x080486c1 : add byte ptr [eax], al ; add cl, cl ; ret
0x08048438 : add byte ptr [eax], al ; add esp, 8 ; pop ebx ; ret
0x080486c2 : add byte ptr [eax], al ; leave ; ret
0x080488a0 : add cl, byte ptr [eax + 0xe] ; adc al, 0x41 ; ret
0x080486c3 : add cl, cl ; ret
0x0804889c : add eax, 0x2300e4e ; dec eax ; push cs ; adc al, 0x41 ; ret
0x080485c5 : add eax, 0x804a064 ; add ecx, ecx ; ret
0x08048582 : add eax, edx ; sar eax, 1 ; jne 0x804858f ; ret
0x080485ca : add ecx, ecx ; ret
```


ROP

```
#include <stdio.h>
#include <stdlib.h>

char *shell = "/bin/sh";

int main(void)
{
    setvbuf(stdout, 0LL, 2, 0LL);
    setvbuf(stdin, 0LL, 1, 0LL);

    char buf[100];

    printf("This time, no system() and NO SHELLCODE!!!\n");
    printf("What do you plan to do?\n");
    gets(buf);

    return 0;
}
```

Reference

- [ROP](#)
- [Mechanisms preventing buffer overflow](#)
- [pwntools packing and unpacking of strings](#)
- [ROPgadget](#)
- [Linux syscall reference](#)

Practice

- train.cs.nctu.edu.tw
 - ret2libc
 - ROP
- secprog.cs.nctu.edu.tw
 - 10/21 practice BOF1
 - 10/21 practice BOF2
 - 10/21 practice BOF3
 - 11/4 practice BOF4
 - 11/4 practice BOF6